

Executive Memo — Collections Recovery Review

Decision

Do not treat the reported 11% MoM improvement as evidence of a structural operational improvement. The clean data shows that February→March cash recovered rose **11.11%**, but the more decision-relevant recovery rate per targeted account rose only **1.38%**, recovery per targeted account rose **1.18%**, and recovery per agent-hour rose **5.38%**. Across complete months January→July, cash recovered was essentially flat (**0.01%**).

What happened?

- The observed dataset covers **1 Jan–8 Aug 2026**, so August is partial and is excluded from complete-month trend conclusions.
- After exact duplicate-event cleaning, February→March successful cash increased from about **■170.3M** to **■189.2M (+11.11%)**. The same increase is visible in recovered accounts (**+11.32%**), so the 11% figure is not fabricated; it is **too narrow to support the broader claim that recovery performance improved 11% MoM**.
- The targeting base also grew **9.81%** from February to March. Consequently, recovery-account rate moved only from **42.11% to 42.69%**.
- Recovery per recovered account fell **0.19%** in March, so higher cash was driven primarily by more recovered accounts, not larger collections per recovered account.

Why?

Fact: The portfolio mix is broadly stable by loan type and risk segment; monthly payment mix swings are modest. Vendor call answer rates range only from about 19.3% to 21.1%, with no clear infrastructure break.

Strong Evidence: Data quality is materially distorting attribution. 98.46% of calls have a borrower_id that disagrees with the account's borrower_id, and the same pattern occurs across other event tables. Borrower-level dimensions should therefore not be trusted.

Strong Evidence: Targeting quality is weak. Of 9,021 targets labeled PROMISE_BROKEN, 8,376 (**92.85%**) have no prior broken PTP event in the available history.

Strong Evidence: Time semantics are unsafe: 30,191 status rows have naive recorded_at earlier than event_at, and ~66.7% of calls disagree with the account timezone. Hour/day analysis should not be used for investment decisions until timestamps are normalized under a documented timezone contract.

Correlation only: 7-day target-to-payment conversion is highest for VOICE (**1.98%**) and lowest for FIELD (**1.67%**), but the assignment data does not support a causal comparison because treatment is not randomized and borrower identity is corrupted.

Is 11% real?

Narrow claim: yes for Feb→Mar cash volume. Broad operational claim: no. A defensible KPI stack should prioritize recovery/account, recovery per targeted account, recovery per agent-hour, and net successful cash rather than one month of cash-volume growth.

■10 Cr recommendation

Area: Better borrower targeting — but stage-gated, not a blind ■10 Cr deployment. This is the only option directly supported by a concrete data defect: targeting rules are demonstrably misaligned with the events they claim to target. In contrast, vendor answer-rate differences are small and there is no clean evidence for incremental agents, AI voice, digital or field expansion.

Financial reality

Using 7-day target-attributed recovery, the observed baseline is approximately **■93.7M annualized** at the observed targeting volume. Scenario incremental recovery from targeting improvement is:

A full **■10Cr** investment would require roughly **107% relative uplift** in target-attributed recovery in one year to break even, which is not supported by the observational data. Therefore the financially correct recommendation is to **fund the targeting capability as a controlled pilot / stage gate**, not to claim that **■10Cr** will generate positive ROI from the current evidence.

Experiment required

Randomize a treatment (new targeting) and control (current targeting) within DPD/risk/loan-type strata; preserve an immutable assignment timestamp; measure **30-day net recovered ■ per eligible account**, plus contact/RPC/PTP/kept-PTP and cost per **■ recovered**. Pre-register the primary metric and keep a 10% holdout until uplift is statistically stable.

Confidence

- **High:** duplicate cleaning, borrower-ID corruption, targeting-rule mismatch, narrow Feb→Mar +11.11% cash result.
- **Medium:** stable portfolio/vendor mix conclusions.
- **Low:** channel/agent causal effects and any ROI estimate, because costs, clean treatment assignment, historical DPD/balance and several requested borrower dimensions are absent.